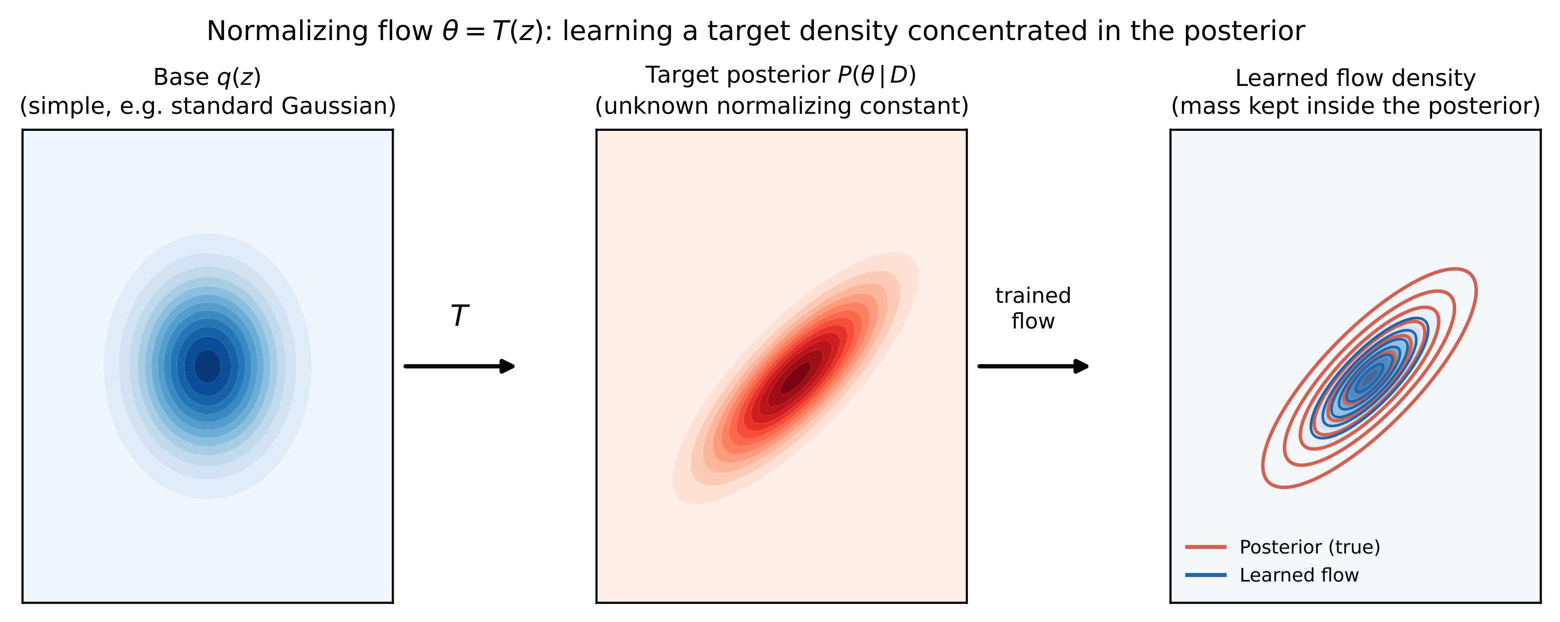
Normalizing flow $\theta = T(z)$: learning a target density concentrated in the posterior
Base $q(z)$
(simple, e.g. standard Gaussian)
T
Target posterior $P(\theta \mid D)$
(unknown normalizing constant)
trained
flow
Learned flow density
(mass kept inside the posterior)
Posterior (true)
Learned flow