


A Bayesian Neural Network: every weight is a distribution, not a number (illustrative diagram)



Each little bell curve is one weight $w_{ij} \sim \mathcal{N}(\mu_{ij}, \sigma_{ij}^2)$

narrow bump = confident (low variance); wide bump = uncertain (high variance).