





Forget-Me-Not: hierarchical mixture over segmentations of $x_{1:n}$



depth 0: whole string, model pool $\mathcal{M}_1$



$\xi(x_{a:b}) = \text{Bayes mixture over pool } \mathcal{M}_a$



each leaf: $p(\cdot | x_{a:b})$ (base measure, adaptive)