

Chapter 9: Four Variants of the Vehicle Routing Problem

Vehicle Routing: Problems, Methods, and Applications
2nd Edition (2014) — Irnich, Schneider & Vigo

Chapter 9: Four Variants of the Vehicle Routing Problem




Table of Contents I

- 1 9.1 Introduction
- 2 9.2 VRP with Backhauls (VRPB)
- 3 9.3 Heterogeneous or Mixed Fleet VRP (HFVRP)
- 4 9.4 Periodic Routing Problems (PVRP)
- 5 9.5 VRP with Split Deliveries (SDVRP)
- 6 9.6 Python Implementations
- 7 9.7 Conclusions and Future Research
- 8 Appendix: Additional Figures

Introduction: Why Study VRP Variants? I

The classical Vehicle Routing Problem (VRP) asks us to find the cheapest set of routes for a fleet of vehicles, all starting and ending at a central depot, such that every customer is visited exactly once and no vehicle exceeds its capacity. In the real world, however, three additional features arise so frequently that they have each generated their own substantial research literature:

Introduction: Why Study VRP Variants? II

Four Variants Covered in Chapter 9

- 1 **VRP with Backhauls (VRPB)** — customers are split into two groups: *linehaul* customers who receive deliveries from the depot, and *backhaul* customers from whom goods must be collected and brought back to the depot. Because many vehicles can only be loaded from the rear, all deliveries must be completed before any pickups begin on the same route.
- 2 **Heterogeneous or Mixed Fleet VRP (HFVRP)** — the company operates vehicles of different sizes, each with its own capacity, fixed ownership cost, and per-kilometre variable cost. Choosing *which type* of vehicle to send is now part of the decision.
- 3 **Periodic VRP (PVRP)** — instead of a single day, planning spans several days (T periods). Each customer has a required *visit frequency* (e.g. twice a week) and the planner must assign which specific days to visit and build efficient daily routes.
- 4 **VRP with Split Deliveries (SDVRP)** — the classical requirement that each customer is served by exactly one vehicle is relaxed. A customer with a large demand may be served by two or more vehicles carrying partial loads, potentially reducing the total number of vehicles needed.

Introduction: Why Study VRP Variants? III

Each of these four families introduces a *characteristic complication* that makes the problem strictly harder than the basic VRP, yet each also opens new algorithmic opportunities. The variants are important both in isolation and in combination: real logistics systems routinely combine, say, periodic scheduling with a heterogeneous fleet.

Key Insight

None of these four variants is a trivial extension of the basic VRP. Each requires specially tailored algorithms for effective solution. The chapter surveys exact methods, classical heuristics, and modern metaheuristics for all four families, together with benchmark results that allow a reader to compare algorithmic performance.

9.2 VRPB: Problem Definition I

In many distribution networks — particularly grocery retail and pharmaceutical distribution — vehicles not only *deliver* goods to customers but also *collect* empty bottles, recyclable packaging, or returned merchandise on the same trip. The two activities create an important asymmetry: once the vehicle has started picking up goods, the cargo area fills from the rear, making it impossible (or at least very inconvenient) to deliver further goods stored deeper in the vehicle.

9.2 VRPB: Problem Definition II

Formal Definition

Let $V = \{0\} \cup L \cup B$ be the vertex set, where:

- vertex 0 is the **depot**,
- $L = \{1, \dots, n_L\}$ is the set of **linehaul** customers (demand $d_i > 0$, meaning goods are *delivered* to them),
- $B = \{n_L + 1, \dots, n_L + n_B\}$ is the set of **backhaul** customers (demand $d_i > 0$, meaning goods are *collected* from them).

Each vehicle has capacity Q . A feasible VRPB solution must satisfy:

- 1 Every linehaul customer in L is visited exactly once.
- 2 Every backhaul customer in B is visited exactly once.
- 3 On every route that serves both linehaul and backhaul customers, **all linehaul visits must come before any backhaul visit** (the precedence constraint).
- 4 The total load on any vehicle (deliveries plus pickups considered separately) never exceeds Q .

9.2 VRPB: Problem Definition III

The Precedence Constraint — Why It Exists

In practice, a vehicle's cargo area is often only accessible from the rear. When the vehicle leaves the depot it is fully loaded with goods for linehaul customers; these are stacked from the front. If the driver tried to pick up backhaul goods first, they would need to rearrange the entire load — an operation that is either physically impossible or so time-consuming as to be uneconomic. The precedence constraint models this physical reality:


if route r visits both linehaul and backhaul customers, then $i \in L$ appears before $j \in B$ in route r .

Key Insight

The VRPB belongs to the broader class of *Delivery and Collection VRP*. Combining both activities in one route is more efficient than running separate delivery and collection tours, but the precedence constraint restricts how much the routes can be interleaved.

VRPB: Three Structural Variants I

The literature distinguishes three ways in which linehaul (L) and backhaul (B) customers can coexist on a route. Relaxing the strict precedence constraint leads to progressively easier (but still NP-hard) problems.



VRPB: Three Structural Variants II

(a) Classical VRPB — Strict Precedence

All linehaul customers on a route are served before any backhaul customer. This corresponds to the physical loading constraint described above. Mixing linehaul and backhaul sequences is simply *forbidden*.

(b) Mixed VRPB (VRPMB) — Interleaving Allowed

The precedence constraint is dropped: a route may alternate freely between linehaul and backhaul customers. This is feasible if the vehicle can be loaded from both ends, or if a forklift is available at each stop. The VRPMB is generally easier to solve because it admits a larger feasible set.

(c) Clustered VRPB (VRPCB) — Grouped Visits

Linehaul and backhaul customers are served in *clusters*: a sub-route visits a group of linehaul customers, then the vehicle may return to the depot (or proceed to a backhaul cluster). This is a compromise between strict and mixed variants.

VRPB: Route Feasibility and Load Tracking I

Understanding how the vehicle load changes along a VRPB route is essential for checking feasibility. Consider a route serving linehaul customers $\ell_1, \ell_2, \dots, \ell_p$ followed by backhaul customers $b_1, b_2, \dots, b_q$.

VRPB: Route Feasibility and Load Tracking II

Load Evolution on a Feasible VRPB Route

Let Q denote the vehicle capacity, d_i the demand of customer i . The vehicle starts at the depot with load $\sum_{j=1}^p d_{\ell_j}$ (carrying all linehaul deliveries). The load *decreases* as linehaul deliveries are made, then *increases* as backhaul pickups are made:

$$\text{Load after serving } \ell_k = \sum_{j=1}^p d_{\ell_j} - \sum_{j=1}^k d_{\ell_j} = \sum_{j=k+1}^p d_{\ell_j}.$$

After the last linehaul customer, the load reaches its minimum. Then:

$$\text{Load after serving } b_k = \sum_{j=1}^k d_{b_j}.$$

Feasibility requires:

$$\sum_{j=1}^p d_{\ell_j} \leq Q \quad (\text{capacity at departure}) \quad \text{and} \quad \sum_{j=1}^q d_{b_j} \leq Q \quad (\text{capacity on return}).$$

VRPB: Mathematical Formulation I

The VRPB can be modelled as an integer program. We define a complete directed graph $G = (V, A)$ where $V = \{0\} \cup L \cup B$ and A is the set of all arcs.

Decision Variables and Parameters

- c_{ij} — travel cost (distance or time) from node i to node j .
- $x_{ij} \in \{0, 1\}$ — equals 1 if some vehicle travels arc (i, j) , and 0 otherwise. (Here x_{ij} is read “x sub i j.”)
- m — number of available vehicles.
- d_i — demand of customer i (delivery amount for $i \in L$, pickup amount for $i \in B$).
- Q — vehicle capacity (maximum load it can carry).
- u_i — auxiliary variable used to eliminate sub-tours.

Objective Function

$$\min \sum_{(i,j) \in A} c_{ij} x_{ij}$$

This formula says: minimise the total travel cost summed over all arcs actually used (i.e. all arcs where $x_{ij} = 1$). Increasing c_{ij} for a long arc makes the model prefer shorter routes.

VRPB: Mathematical Formulation III

Constraints

Degree constraints (every customer visited exactly once):

$$\sum_{j \in V} x_{ij} = 1 \quad \forall i \in L \cup B \quad (\text{each customer has exactly one outgoing arc})$$

$$\sum_{i \in V} x_{ij} = 1 \quad \forall j \in L \cup B \quad (\text{each customer has exactly one incoming arc})$$

Fleet size (depot degree):

$$\sum_{j \in L} x_{0j} \leq m \quad (\text{at most } m \text{ routes depart from the depot})$$

Precedence constraint (linehaul before backhaul): For every arc (i, j) with $i \in B$ and $j \in L$, we set $x_{ij} = 0$. Equivalently, no arc may go from a backhaul node directly to a linehaul node.

Capacity constraints: Using sub-tour elimination variables u_i (MTZ formulation):

$$u_i - u_j + Q \cdot x_{ij} \leq Q - d_j \quad \forall i, j \in V, i \neq j$$

VRPB: Exact Methods I

Exact methods for VRPB guarantee finding the globally optimal solution but are computationally expensive. The two main approaches are **Branch-and-Bound** and **Column Generation** (Branch-and-Price).

Branch-and-Bound for VRPB

The Branch-and-Bound (B&B) approach works as follows. Start by solving the *Linear Programming relaxation* (LP) of the ILP — this gives a lower bound on the optimal integer solution. If the LP solution contains fractional x_{ij} values, we *branch*: fix one fractional variable to 0 or 1 and solve two sub-problems. This creates a binary search tree; at each node we compute bounds and prune branches that cannot improve the current best integer solution.

For VRPB, the first exact approach was proposed by **Goetschalckx and Jacobs-Blecha (1989)**, who used a set-partitioning formulation. The method can solve instances with up to 70 customers to optimality.

Branch-and-Price (Column Generation)

A more powerful exact approach decomposes the problem into:

- **Master problem:** select a subset of pre-generated routes (columns) with minimum total cost such that every customer is covered.
- **Pricing sub-problem:** find new routes (columns) with negative reduced cost that should be added to the master problem. For VRPB, this is a *constrained shortest path problem* that respects the precedence constraint.

Toth and Vigo extended B&P to handle both VRP and VRPB. Using **Valid Inequalities** (capacity cuts, precedence cuts) dramatically tightens the LP bound.

VRPB: Exact Methods III

Computational Complexity

The VRPB is NP-hard (it contains VRP as a special case when $B = \emptyset$). In practice:

- Exact B&B can handle $|L| + |B| \lesssim 50$ customers.
- Branch-and-Price pushes this to $|L| + |B| \lesssim 100$ customers.
- Beyond this, heuristic methods are essential.

Benchmark: Toth and Vigo (1997) Results

On the standard VRPB benchmark set (Goetschalckx–Jacobs–Blecha instances, $n \leq 100$ customers):

- B&P solves most instances to optimality within minutes.
- Instances with $n = 100$ may require several hours.
- The LP lower bound is typically within 3–5% of the optimum.

VRPB: Heuristics I

Because exact methods are too slow for large instances, a rich collection of heuristics has been developed for VRPB. The most important ones are adaptations of classical VRP heuristics that respect the precedence constraint.

Clarke-Wright Savings Heuristic (Adapted for VRPB)

The **savings algorithm** (Clarke and Wright, 1964) starts with a “star” solution in which each customer has its own dedicated route: Depot $\rightarrow$ customer $i \rightarrow$ Depot. It then merges pairs of routes by computing the *saving* $s(i, j) = c_{0i} + c_{j0} - c_{ij}$: the distance saved by linking customer i to customer j instead of returning to the depot between them.

The saving says: instead of travelling Depot $\rightarrow i \rightarrow$ Depot $\rightarrow j \rightarrow$ Depot (cost $c_{0i} + c_{i0} + c_{0j} + c_{j0}$), travel Depot $\rightarrow i \rightarrow j \rightarrow$ Depot (cost $c_{0i} + c_{ij} + c_{j0}$). The saving is $c_{i0} + c_{0j} - c_{ij}$.

For VRPB, the savings algorithm must additionally ensure:

- 1 Merging two routes never creates a backhaul-before-linehaul sequence.
- 2 The combined load does not exceed Q .

In practice, this means linehaul customers are merged first, then backhaul customers, and only feasible merges are accepted.

VRPB: Heuristics III

Worked Example: VRPB Savings Algorithm

Consider depot $D = (0, 0)$ and 4 customers with Euclidean distances:

	D	$L1$	$L2$	$B1$	$B2$
D	0	3.6	2.8	2.2	3.2
$L1$	3.6	0	2.0	4.5	5.0
$L2$	2.8	2.0	0	3.1	4.2
$B1$	2.2	4.5	3.1	0	1.4
$B2$	3.2	5.0	4.2	1.4	0

Vehicle capacity $Q = 10$; demands: $L1 = 4, L2 = 5, B1 = 3, B2 = 4$.

Step 1. Initial solution: 4 separate routes, total cost $= 2(3.6 + 2.8 + 2.2 + 3.2) = 23.6$.

Step 2. Compute savings for linehaul pairs:

$$s(L1, L2) = 3.6 + 2.8 - 2.0 = 4.4 \quad (\text{best!})$$

Merge $L1$ and $L2$: route $D \rightarrow L1 \rightarrow L2 \rightarrow D$, load $= 9 \leq 10$.

VRPB: Metaheuristics I

For large instances (hundreds of customers), metaheuristics that intelligently explore the solution space are the method of choice.

Tabu Search for VRPB

Tabu Search (TS) is a local search method that accepts moves even when they temporarily worsen the solution, to escape local optima. It maintains a *tabu list* of recently visited solutions (or recently made moves) and forbids revisiting them for a fixed number of iterations.

For VRPB, Tabu Search operators include:

- **Customer relocation:** move a single customer from one route to another (checking precedence and capacity).
- **2-opt:** reverse a segment of a route. Only applicable if the resulting route still has all linehaul before all backhaul.
- **Or-opt:** move a chain of 1, 2, or 3 consecutive customers to another position in any route.
- **Cross-exchange:** swap segments between two routes.

The most effective published TS for VRPB is due to **Toth and Vigo (1997)**, achieving solutions within 1–2% of the best known on benchmark instances.

Genetic Algorithms (GA) for VRPB

A Genetic Algorithm maintains a *population* of solutions (“chromosomes”). At each generation it selects parent solutions by fitness (total cost), applies *crossover* to combine parents into offspring, and applies *mutation* to introduce new material. For VRPB:

- **Chromosome encoding:** a permutation of customers with position-based encoding; backhaul customers are always placed after linehaul customers in each route segment.
- **Crossover:** Order Crossover (OX) preserves the relative order of customers from one parent while copying a segment from the other parent.
- **Repair operator:** after crossover, a feasibility check ensures the precedence constraint holds; violations are fixed by moving misplaced customers.

Validating benchmark results, GAs reach solutions of comparable quality to TS in similar or shorter times, but are more sensitive to parameter tuning.


Adaptive Large Neighbourhood Search (ALNS)

ALNS (Ropke and Pisinger, 2006) is among the most competitive methods for VRP variants. It repeatedly *destroys* part of the current solution (removing a random subset of customers) and *repairs* it (re-inserting removed customers optimally). Destroy and repair operators are selected adaptively: operators that have historically produced good solutions are chosen more frequently. For VRPB, the repair step must respect the precedence constraint when re-inserting backhaul customers.


VRPB: Route Diagram and Summary I

VRP with Backhauls (VRPB): Precedence Constraint

Infeasible: Backhaul served before Linehaul



Feasible: All Linehauls First, then Backhauls



VRPB: Route Diagram and Summary II

The left panel shows an *infeasible* route where a backhaul customer is served before linehaul customers — this violates the precedence constraint. The right panel shows two feasible routes: Route 1 (blue) serves linehaul customers $L1, L2, L4$ before backhaul customers $B1, B3$; Route 2 (orange) serves $L3$ before $B2$.

VRPB Summary

The VRP with Backhauls models the practical reality that vehicles often both deliver and collect on the same trip. The precedence constraint (linehaul before backhaul) reflects vehicle loading geometry. Exact methods using B&P solve instances up to ≈ 100 customers. For larger instances, metaheuristics such as Tabu Search and ALNS achieve solutions within 1–3% of the best known.

9.3 HFVRP: Problem Definition and Variants I

In the classical VRP, all vehicles are assumed identical: same capacity, same fuel consumption, same fixed cost. In practice, almost every real fleet is *heterogeneous*: a logistics company might operate motorcycles for small express deliveries alongside 40-tonne articulated lorries for bulk shipments. The Heterogeneous Fleet VRP (HFVRP) — also called Mixed Fleet VRP or Free Fleet Size and Mix VRP — explicitly models this diversity.

Problem Setting

There are T vehicle *types*, indexed $t = 1, \dots, T$. Each type t is characterised by:

- Q_t — capacity (maximum load it can carry per trip),
- f_t — fixed cost per vehicle used (ownership or daily rental cost),
- r_t — variable cost per unit of distance travelled.

There may be m_t vehicles of type t available (fixed fleet size case), or the fleet size may be unlimited (free fleet size case, FSMVRP). The goal is to assign vehicle types to routes and find routes that minimise total cost (fixed + variable).

9.3 HFVRP: Problem Definition and Variants II

HFVRP Sub-Families (Table 9.1 in the Book)

HFVRP Problem Variants (from Table 9.1 in the book)

Acronym	Fleet Size	Vehicle Types	Fixed Cost	Variable Cost
HFVRP	Fixed	Heterogeneous	Yes	Yes
FSMVRP	Free	Heterogeneous	Yes	Yes
FSMF	Free	Heterogeneous	Yes	No
FSMV	Free	Heterogeneous	No	Yes
MDVRP	Fixed	Homogeneous	No	Yes

The key distinction is between **fixed fleet** (HFVRP: you must use exactly the vehicles you have) and **free fleet** (FSMVRP: you choose how many of each type to use). In the free-fleet case, using many small vehicles trades lower capacity per vehicle for lower fixed costs (if small vehicles are cheaper to own).

HFVRP: Mathematical Formulation I

Let $G = (V, A)$ with $V = \{0, 1, \dots, n\}$ ($0 = \text{depot}, 1, \dots, n = \text{customers}$). Let K_t be the index set of vehicles of type t .

Variables

- $x_{ij}^{kt} \in \{0, 1\}$ — equals 1 if vehicle k of type t traverses arc (i, j) . Read as “x sub ij super kt.”
- $y_k^t \in \{0, 1\}$ — equals 1 if vehicle k of type t is used (i.e. its route is non-empty).
- d_i — demand of customer i .

HFVRP: Mathematical Formulation II

Objective Function

$$\min \sum_t \left[f_t \sum_k y_k^t + r_t \sum_k \sum_{(i,j) \in A} c_{ij} x_{ij}^{kt} \right]$$

Here:

- $f_t \sum_k y_k^t$ — total fixed cost for all vehicles of type t that are actually used.
- $r_t \sum_k \sum_{(i,j)} c_{ij} x_{ij}^{kt}$ — total variable (distance-based) cost for all routes driven by vehicles of type t .
- c_{ij} — distance or travel time from node i to node j .

The formula is balancing two opposing forces: using *fewer, larger* vehicles reduces fixed costs but may require them to travel longer routes; using *more, smaller* vehicles may give shorter routes but increases fixed costs.

HFVRP: Mathematical Formulation III

Constraints

Each customer visited exactly once:

$$\sum_t \sum_k \sum_{j \in V} x_{ij}^{kt} = 1 \quad \forall i \in V \setminus \{0\}$$

Capacity constraint for each vehicle:

$$\sum_{i \in V} d_i \sum_{j \in V} x_{ij}^{kt} \leq Q_t \cdot y_k^t \quad \forall k, \forall t$$

This says: the total demand served by vehicle k of type t cannot exceed Q_t (the capacity of type t). If $y_k^t = 0$ (vehicle not used), the right-hand side is 0 and the vehicle serves no customers.

Fleet size constraint:

$$\sum_k y_k^t \leq m_t \quad \forall t \quad (\text{at most } m_t \text{ vehicles of type } t)$$


For the free-fleet variant (FSMVRP), this constraint is dropped.

Flow conservation and sub tour elimination: standard formulations as for VRP (e.g. MTZ constraints)

HFVRP: Example with Multiple Vehicle Types I

HFVRP: Example with Multiple Vehicle Types II

Heterogeneous Fleet VRP (HFVRP): Different Vehicle Types



HFVRP: Exact Methods II

Set-Partitioning Formulation (Column Generation)

The standard exact approach for HFVRP uses a **set-partitioning** formulation where each column corresponds to a complete route for a specific vehicle type. Let $\mathcal{R}_t$ be the set of all feasible routes for vehicle type t , and let $a_{ir} = 1$ if customer i appears in route r . The formulation is:

$$\min \sum_t \sum_{r \in \mathcal{R}_t} (f_t + \text{cost}_r) \cdot \lambda_r$$

subject to:

$$\sum_t \sum_{r \in \mathcal{R}_t} a_{ir} \lambda_r = 1 \quad \forall i \quad (\text{each customer covered exactly once})$$

$$\lambda_r \in \{0, 1\} \quad \forall r \quad (\text{select or don't select each route})$$

where λ_r (lambda, here λ is read as “lambda sub r ”) is 1 if route r is selected and 0 otherwise, and f_t is the fixed cost of the vehicle type used for route r .

The LP relaxation is solved by Column Generation; the pricing sub-problem finds the route of minimum reduced cost for each vehicle type.

Clarke-Wright Savings (Adapted)

The savings algorithm is adapted for HFVRP by choosing the cheapest vehicle type for each route. After merging two routes, the algorithm selects the smallest vehicle type that has sufficient capacity for the combined route. This is efficient but may be suboptimal because it does not globally optimise vehicle-type assignments.

Tabu Search (Gendreau, Laporte, Musaraganyi, Taillard)

The most influential metaheuristic for HFVRP was proposed by **Gendreau et al.** (1999). Their Tabu Search uses:

- 1 **GENIUS insertion**: a powerful initial solution constructor that generalises nearest-neighbour insertion.
- 2 **Move types**: single customer relocation, 2-opt, 3-opt, and vehicle type swap (change the vehicle assigned to an entire route).
- 3 **Intensification**: restart from the best solution found and search more deeply.
- 4 **Diversification**: when the search stagnates, randomly perturb the current solution.

On the 12 Golden instances, this TS produces solutions within 2–4% of the best known, with runtimes of seconds to minutes.

HFVRP: Heuristics and Metaheuristics III

Genetic Algorithm (Prins, 2009)

Prins (2009) proposed a GA with a chromosome encoding that simultaneously represents the permutation of customers and the vehicle-type assignment. The crossover operator (LAX) preserves good partial routes from each parent. A local search phase (2-opt, Or-opt) refines each offspring before it enters the population. This hybrid GA is competitive with the best TS results on Golden instances.

Simulated Annealing (Imran et al., 2009)


Simulated Annealing (SA) — named after the annealing of metals, where slow cooling allows atoms to find a low-energy state — accepts *worse* solutions with a probability that decreases over time:

$$P(\text{accept worse}) = \exp(-\Delta\text{cost} / T_{\text{current}})$$

where T_{current} (the “temperature”) starts high (accepting most moves) and gradually decreases (becoming more selective). **Imran et al. (2009)** apply SA with a perturbation operator that swaps vehicle types between routes, achieving results competitive with Tabu Search on HFVRP benchmarks.

HFVRP: Heuristics and Metaheuristics IV

Performance Comparison



The left panel shows approximate percentage gaps from the best known solutions on Golden instances. Modern metaheuristics (TS, GA, SA) all achieve gaps of 1–3%.

9.4 PVRP: Problem Definition I

In many real logistics settings, routing is not a one-time decision but a *recurring* one. A pharmaceutical distributor visits hospitals every day; a bakery delivery service visits restaurants three times a week; a waste collection company empties bins once or twice a week. The **Periodic Vehicle Routing Problem (PVRP)** captures this multi-period planning structure.

9.4 PVRP: Problem Definition II

Formal Definition

The planning horizon consists of T periods (days). There are n customers; each customer i has:

- $f_i \in \{1, 2, \dots, T\}$ — **visit frequency**: the number of times customer i must be visited over the T -period horizon.
- A set $\mathcal{P}_i$ of **admissible visit patterns**: binary vectors $p = (p_1, \dots, p_T)$ of length T with exactly f_i ones, indicating on which days the customer is visited. For example, with $T = 5$ and $f_i = 2$, a valid pattern is $(1, 0, 0, 1, 0)$ (Monday and Thursday).
- d_i — demand per visit (often assumed constant across visits).

The PVRP seeks to:

- 1 Assign an admissible visit pattern $p_i \in \mathcal{P}_i$ to each customer i .
- 2 For each day $t = 1, \dots, T$, build a set of routes serving all customers i with $p_{it} = 1$.
- 3 Minimise the total routing cost over all days, and optionally balance the daily workload.

9.4 PVRP: Problem Definition III


Why PVRP Is Harder than VRP

The PVRP is strictly harder than VRP for two reasons:

- ① **Pattern assignment:** even before routing, we must decide *which* f_i days to visit each customer. This is itself a combinatorial problem with $|\mathcal{P}_i|$ choices per customer. For $T = 5$ and $f_i = 2$, there are $\binom{5}{2} = 10$ possible patterns.
- ② **Daily routing:** once patterns are assigned, we must solve a *separate* VRP for each of the T days — so we are solving T coupled VRPs simultaneously.

The two sub-problems interact: a good pattern assignment leads to balanced daily demands that are easier to route efficiently.

9.4 PVRP: Problem Definition IV



PVRP: Visit Patterns and Admissibility I

The set of admissible visit patterns for a customer with frequency f_i over a T -period horizon consists of all binary vectors of length T with exactly f_i ones. However, not all such vectors are *practically admissible*: for example, a customer requiring 3 visits per 5-day week should not be visited three times on consecutive days if their inventory can hold a week's worth of stock. The admissible set $\mathcal{P}_i$ is specified externally based on operational constraints.

PVRP: Visit Patterns and Admissibility II

PVRP: Allowable Visit Patterns for T=5 Day Planning Horizon

f=1 (once/period):

(1,0,0,0,0), (0,1,0,0,0), (0,0,1,0,0), (0,0,0,1,0), (0,0,0,0,1)

f=2 (twice/period):

(1,0,1,0,0), (1,0,0,1,0), (1,0,0,0,1), (0,1,0,1,0), (0,1,0,0,1) ... (6 total)

f=3 (three/period):

(1,1,1,0,0), (1,1,0,1,0), (1,1,0,0,1), (1,0,1,1,0), (1,0,1,0,1) ... (10 total)

f=5 (daily):

(1,1,1,1,1)


PVRP: Visit Patterns and Admissibility III

Standard Benchmark Patterns (Cordeau et al., 1997)

The most widely used PVRP benchmark (“p” instances) uses $T = 5$ and frequencies $f_i \in \{1, 2, 3, 5\}$. The admissible patterns are the standard patterns listed in the figure above. A frequency-5 customer must be visited every day; a frequency-1 customer is visited once in any of the 5 days. The assignment of patterns determines which customers appear in each day’s routing problem.

PVRP: Multi-Day Routes I

PVRP: Different Routes Each Day (T=3 day excerpt)



PVRP: Multi-Day Routes II

This figure shows how the same set of customers generates different daily routes depending on their visit patterns. On Day 1 (left), customers $C1$, $C2$, $C5$ (high-frequency) and $C3$, $C6$ (who happen to have Day 1 in their pattern) are visited. On Day 2 (middle), a different subset is active. The key point is that *route structures change day to day*, and the planner must design all T days' routes simultaneously.

Workload Balancing

A secondary objective often included in PVRP is to make daily workloads as equal as possible. If all high-demand customers are assigned to Monday, Monday's routing problem is large and expensive while other days are easy. Good pattern assignment distributes customers evenly over the week.

PVRP: The Standard (Benchmark) Formulation I

The standard PVRP formulation, introduced by **Cordeau, Gendreau, and Laporte (1997)**, is a three-index integer program.

Parameters and Variables

- n — number of customers.
- T — number of periods (days) in the planning horizon.
- m — number of vehicles available per day.
- Q — vehicle capacity.
- d_i — demand of customer i per visit.
- c_{ij} — travel cost between nodes i and j .
- $\mathcal{P}_i$ — set of admissible visit patterns for customer i .
- z_{ip} — equals 1 if pattern $p \in \mathcal{P}_i$ is selected for customer i , and 0 otherwise. (“ z sub i p ”).
- x_{ij}^{tk} — equals 1 if vehicle k traverses arc (i,j) on day t .

PVRP: The Standard (Benchmark) Formulation II

Objective

$$\min \sum_{t=1}^T \sum_{k=1}^m \sum_{(i,j) \in A} c_{ij} x_{ij}^{tk}$$

This sums travel cost over all days $t = 1, \dots, T$, all vehicles $k = 1, \dots, m$, and all arcs used.

PVRP: The Standard (Benchmark) Formulation III

Key Constraints

Pattern selection (exactly one pattern per customer):

$$\sum_{p \in \mathcal{P}_i} z_{ip} = 1 \quad \forall i \in \{1, \dots, n\}$$

Linking patterns to daily routes: if customer i is assigned pattern p with $p_t = 1$ (visited on day t), then i must appear in day t 's routing:

$$\sum_{j \in V} x_{ij}^{tk} = \sum_{p: p_t=1} z_{ip} \quad \forall i, \forall t$$

Daily VRP constraints: for each day t , the routes $\{x^{tk}\}$ form a feasible VRP solution: flow conservation, capacity, and sub-tour elimination.

Fleet size: at most m routes per day.

Exact Methods

Exact methods for PVRP are based on Branch-and-Price. The sub-problems decompose by day: for each day t , find routes of minimum reduced cost that cover the customers assigned to day t . Cordeau, Gendreau, and Laporte (1997) applied this approach and solved the benchmark instances with up to 101 customers. The LP lower bound is typically within 2–5% of the optimum.

Tabu Search (Cordeau et al., 1997 — the definitive standard)

The most influential heuristic for PVRP is the **Unified Tabu Search** of Cordeau, Gendreau, and Laporte (1997), which handles multiple VRP variants under one framework. The algorithm:

- 1 Starts with a greedy initial solution (insert customers into days that minimise cost increase, choosing among admissible patterns).
- 2 At each iteration, moves a customer from one day's route to another (or changes their pattern assignment) by choosing the move with minimum cost increase (or maximum improvement).
- 3 Forbids recently performed moves (tabu list) to prevent cycling.
- 4 Applies penalties for constraint violations to allow temporary infeasibility and escape local optima.

This TS routinely matches or improves the best known solutions on the standard PVRP benchmark.

Metaheuristic Variants for PVRP

Several extensions of the standard TS have been proposed:

- **Granular Tabu Search** (Toth and Vigo, 2003): restricts moves to “promising” arcs (short arcs), dramatically reducing the neighbourhood size without sacrificing solution quality. The *granularity parameter* β controls how many arcs are considered; typical values are $\beta = 1.5\text{--}3.0$.
- **Variable Neighbourhood Search (VNS)**: systematically explores different neighbourhood structures. For PVRP, the neighbourhoods include pattern reassignment, within-day 2-opt, and cross-day customer relocation.
- **Evolutionary Local Search (ELS)**: combines a population of solutions with local search. Each solution in the population is locally optimised; new solutions are generated by combining parts of high-quality solutions.

PVRP with Time Windows (PVRPTW)

A natural extension adds *time windows* to the PVRP: each customer must be visited within a specified time interval on each visit day. This makes the sub-problem for each day a *VRPTW*, which is itself harder than VRP. Heuristics for PVRPTW typically use the same TS framework but with a more complex feasibility check that includes time-window violations.

PVRP: Other Variants I

Multi-Depot PVRP (MDPVRP)

When there are multiple depots (each with its own fleet), the pattern assignment problem must also decide *which depot* serves each customer on each visit day. This adds a third dimension to the decision space. Valid approaches extend the PVRP TS by adding a depot-assignment operator: move a customer between depots as well as between days.

PVRP with Service Time Dependence

In some applications, the service time or demand at a customer depends on the day of visit (e.g. supermarkets order more on Fridays before weekends). This time-dependent variant requires the pattern assignment to account for varying per-visit demands.

PVRP as a Two-Phase Problem

A popular decomposition approach solves PVRP in two phases:

- 1 **Phase 1 (Pattern Assignment)**: assign a visit pattern to each customer, aiming to balance the daily customer load. This can be solved approximately by a bin-packing or set-covering heuristic.
- 2 **Phase 2 (Daily Routing)**: given the pattern assignments, solve T independent VRPs (one per day) using a standard VRP solver.

Two-phase methods are fast and easy to implement but may sacrifice quality because Phase 1 does not consider routing costs. Integrated methods that solve both phases simultaneously generally outperform them.

9.5 SDVRP: Problem Definition I

In the classical VRP, each customer must be served by *exactly one vehicle*. This constraint can be wasteful: if a customer requires 9 units and the vehicle capacity is 10, then that vehicle is nearly full but carries only one customer's goods. The **Split Delivery VRP (SDVRP)** relaxes this constraint, allowing a customer's demand to be split across multiple vehicles.

Formal Definition

We have n customers, each with demand $d_i > 0$, and m vehicles each with capacity Q . The SDVRP allows customer i to be visited by multiple vehicles, receiving a partial delivery $d_i^v > 0$ from vehicle v (where v ranges over all vehicles that visit i), as long as the total delivery equals the full demand:

$$\sum_{v: v \text{ visits } i} d_i^v = d_i \quad \forall i.$$

The **number of times** a customer is visited (split deliveries) is unrestricted (in principle), but in practice most algorithms limit splits to avoid excessive customer inconvenience.

9.5 SDVRP: Problem Definition II


Intuition: Why Splitting Helps

Imagine two customers each needing 8 units, with $Q = 10$. Without splitting: 2 separate routes needed (8 units each, capacity wasted). With splitting: one vehicle carries 8 for customer 1 and 2 for customer 2 (full load); another vehicle carries 6 for customer 2 (also full). We have saved one depot-return trip at the cost of visiting customer 2 twice.


SDVRP: Split vs. No-Split Comparison I

Split Delivery VRP (SDVRP): Customer Served by Multiple Vehicles

Standard VRP: C1 must have dedicated route ($d=8 > Q=5$)



SDVRP: C1 split → Vehicle A delivers 5, Vehicle B delivers 3



SDVRP: Split vs. No-Split Comparison II

The left panel shows the standard VRP where customer $C1$ (demand $8 > Q = 5$) must have a dedicated route. The right panel shows the SDVRP where $C1$ is split: Vehicle A delivers 5 units to $C1$ on one trip, Vehicle B delivers the remaining 3 units on another trip while also serving $C2$ and $C3$. The total number of routes is reduced from 2 to 2 here, but the key saving is that Vehicle A can now also serve $C3$ on its trip — reducing total distance.

The Fundamental Trade-Off in SDVRP

Splitting a customer reduces the number of routes (fewer vehicles needed) but may increase the number of visits to some customers, which is inconvenient for the customer and increases operational complexity. The SDVRP finds the sweet spot: split only when the distance/cost saving justifies the extra visit.

SDVRP: Mathematical Formulation and Basic Variant I

The SDVRP was first studied by **Dror and Trudeau (1989, 1990)**. They showed that the problem on a tree network can be solved in polynomial time, but the general (Euclidean or graph) case is NP-hard.

SDVRP: Mathematical Formulation and Basic Variant II

Basic SDVRP Formulation

Let $x_{ij}^v \in \{0, 1\}$ indicate whether vehicle v traverses arc (i, j) , and let $q_i^v \geq 0$ be the amount delivered to customer i by vehicle v . Then:

$$\min \sum_v \sum_{(i,j) \in A} c_{ij} x_{ij}^v$$

subject to:

$$\sum_v q_i^v = d_i \quad \forall i \quad (\text{full demand satisfied})$$

$$q_i^v \leq d_i \cdot \sum_{j \in V} x_{ij}^v \quad \forall i, v \quad (\text{delivery only if vehicle visits})$$

$$\sum_i q_i^v \leq Q \quad \forall v \quad (\text{capacity per vehicle})$$

Flow conservation and sub-tour elimination constraints complete the model.

SDVRP: Mathematical Formulation and Basic Variant III

Lower Bound: Number of Vehicles Required

A fundamental property of the SDVRP is that the minimum number of vehicles required to serve all customers is:

$$m^* = \left\lceil \frac{\sum_{i=1}^n d_i}{Q} \right\rceil$$

where $\lceil \cdot \rceil$ (the ceiling function) rounds up to the nearest integer. Here $\sum_{i=1}^n d_i$ is the **total demand** (sum of all customers' demands) and Q is the vehicle capacity. This bound follows from the fact that with split deliveries, we can *always* fill each vehicle completely (except possibly the last one), so the minimum number of vehicle-loads needed is exactly $\lceil D/Q \rceil$ where $D = \sum_i d_i$.

In contrast, without splitting, bin-packing constraints may force us to use *more* than $\lceil D/Q \rceil$ vehicles because demands may not pack perfectly.

Worked Example: Savings from Splitting

Five customers with demands $[8, 7, 6, 9, 5]$, total $D = 35$, capacity $Q = 10$.

- **Lower bound** (split allowed): $\lceil 35/10 \rceil = 4$ vehicles minimum.
- **Without splitting** (bin packing): demands $[9, 8, 7, 6, 5]$ sorted descending. First route: 9 only (room for 1 more: $9 + 8 > 10$ so 9 alone); Second: 8 alone; Third: $7 + 3 \leq 10$ but no customer with demand 3 exists; and so on. Result: 5 vehicles needed.
- **With splitting**: 4 vehicles (saving 1 vehicle trip).

SDVRP: Theoretical Properties I

Dror and Trudeau (1989, 1990) proved several key properties that guide the design of algorithms for the SDVRP.

Property 1: Optimal Split Structure

In an optimal SDVRP solution, if a customer i is visited by more than one vehicle, it is visited at most twice (under certain triangle-inequality conditions on the distances). This is a remarkable result: even though in principle a customer could be visited many times, twice is almost always sufficient in practice.

SDVRP: Theoretical Properties II


Property 2: Maximum Savings from Splitting

The cost ratio between the optimal VRP solution (no split) and the optimal SDVRP solution (split allowed) is bounded:

$$\frac{\text{VRP}^*}{\text{SDVRP}^*} \leq 1 + \frac{1}{2\lceil Q/d_{\min} \rceil - 1}$$

where $d_{\min}$ is the smallest customer demand. Here $Q/d_{\min}$ is the ratio of vehicle capacity to the smallest demand; $\lceil Q/d_{\min} \rceil$ is the maximum number of customers a vehicle can carry if all have minimum demand. When this number is large (small demands relative to capacity), the bound approaches 1 — splitting gives little benefit because vehicles are already nearly full. When demands are close to Q , splitting can give savings of up to 50%.

SDVRP: Theoretical Properties III



SDVRP: Exact and Heuristic Methods I

Exact Methods (Archetti, Speranza, Hertz)

Exact algorithms for the SDVRP use Branch-and-Price. The pricing sub-problem is a *shortest path problem with resource constraints*: find a minimum-cost route (path from depot through a subset of customers back to depot) that respects capacity, where customers may be visited partially (split loads). This sub-problem is NP-hard in general, but efficient label-setting algorithms handle it in practice.

Archetti, Speranza, and Hertz (2006) developed a B&P that can solve instances with up to 100 customers to optimality. The key observation is that in the LP relaxation of the set-partitioning master problem, customers with demand $d_i > Q/2$ *must* be split (they cannot fit in a single vehicle with anyone else), so these splits are pre-processed.

Heuristics: Nearest-Neighbour and Insertion

Simple constructive heuristics for SDVRP work as follows:

- 1 Initialise each vehicle at the depot with a full load of Q units.
- 2 At each step, the current vehicle visits the nearest unserved customer (or the nearest customer with remaining undelivered demand).
- 3 If the vehicle can fully serve the customer, do so. If not (demand exceeds remaining capacity), deliver the remaining capacity as a partial delivery, mark the customer as partially served, and return to depot.
- 4 Continue until all customers are fully served.

This greedy approach is fast but may give solutions 10–20% above optimal.

SDVRP: Exact and Heuristic Methods III

Tabu Search for SDVRP (Archetti, Hertz, Speranza, 2006)

The most effective published TS for SDVRP uses these neighbourhoods:

- **Split/unsplit:** take a partially-served customer and either merge the two partial deliveries into one (unsplit) or further split an existing single delivery (create a new partial delivery on another vehicle).
- **Customer relocation:** move a (partial) delivery from one route to another.
- **2-opt within a route:** reverse a segment of a route, check that capacity is not violated.
- **Interchange:** swap deliveries between two routes.

On benchmark instances, this TS produces solutions within 0.5–2% of the best known in seconds.

The SDVRP with Time Windows (SDVRPTW)

Combining split deliveries with time windows is particularly challenging: a split customer must be visited at *different times* by different vehicles, so the time windows must be consistent across all visits. The B&P approach extends to SDVRPTW with additional time-window constraints in the pricing sub-problem.

SDVRP: Variants and Extensions I

SDVRP Variants Overview

VRP variant	Demand type	Fleet size	Notes
VRP (classical)	Delivery only	Fixed, homogeneous	No splits allowed
SDVRP	Delivery only	Fixed, homogeneous	Splits allowed
SDVRPTW	Delivery + time windows	Fixed, homogeneous	Splits + windows
SDVRPB	Delivery + backhaul	Fixed, homogeneous	Splits + backhauls
HFVRP+SD	Delivery only	Free, heterogeneous	Splits + mixed fleet

SDVRP: Variants and Extensions II

SDVRP with Backhauls (SDVRPB)

When split deliveries are combined with backhaul customers, the problem becomes significantly more complex. Now linehaul customers may have their demand split across multiple vehicles, and backhaul pickups must still come after all linehaul deliveries on each route. Heuristics adapt the standard SDVRP operators by checking the precedence constraint after every move.

Savings from Splitting: Empirical Observations

Empirical studies on real-world instances suggest:

- For instances where $d_i \leq Q/3$ for all customers (small demands relative to capacity), splitting gives negligible savings.
- For instances where $d_i \approx Q/2$ for many customers (medium demands), savings of 5–15% in total distance are typical.
- For instances where some customers have $d_i > Q$ (demands larger than capacity), splitting is *mandatory*: the customer *must* be split.

Python: VRPB Savings Algorithm

The following Python code implements the Clarke-Wright savings algorithm adapted for VRPB. It starts with individual routes, computes savings for all linehaul pairs, merges the best feasible pair, then does the same for backhauls, and finally attempts to join linehaul and backhaul routes when capacity allows.

```
import numpy as np
from itertools import combinations

def vrp_b_savings(depot, linehauls, backhauls, demands, Q):
    """
    Clarke-Wright savings algorithm for VRPB.
    depot : (x,y) tuple
    linehauls : dict {name: (x,y)}
    backhauls : dict {name: (x,y)}
    demands : dict {name: demand}
    Q : vehicle capacity
    Returns list of routes, each a list of customer names.
    """
    all_c = {**linehauls, **backhauls}
    # Distance helper
    def dist(a, b):
        return np.hypot(a[0]-b[0], a[1]-b[1])

    # Initial solution: each customer alone
    routes = {c: [c] for c in all_c}
    loads = {c: demands[c] for c in all_c}

    # Compute all savings  $s(i,j) = d(\text{depot}, i) + d(\text{depot}, j) - d(i,j)$ 
```

Python: SDVRP Minimum Vehicles and Greedy Split

This code computes the theoretical minimum number of vehicles for the SDVRP and implements a simple greedy split delivery heuristic.

```
import math

def sdvrp_min_vehicles(demands, Q):
    """
    Compute the lower bound on number of vehicles needed with split deliveries.
    demands : list of customer demands
    Q       : vehicle capacity
    """
    total_demand = sum(demands)
    return math.ceil(total_demand / Q)

def sdvrp_greedy(customers, demands, Q):
    """
    Greedy split delivery heuristic.
    customers : list of (x,y) customer positions (index = customer ID)
    demands   : list of demands (same indexing)
    Q         : vehicle capacity
    Returns list of vehicle routes: each route is list of (customer_id, qty_delivered).
    """
    import numpy as np
    remaining = list(demands)  # remaining demand for each customer
    n = len(customers)
    routes = []
```

Python: PVRP Pattern Assignment

The following code demonstrates pattern assignment for the PVRP: given customer frequencies and admissible patterns, it assigns patterns to balance the daily workload (number of customers per day).

```
from itertools import combinations

def generate_patterns(T, f):
    """
    Generate all binary vectors of length T with exactly f ones.
    T : planning horizon (number of days)
    f : visit frequency
    Returns list of tuples, each a binary pattern.
    """
    return [tuple(1 if i in combo else 0 for i in range(T))
            for combo in combinations(range(T), f)]

def pvrp_balance_assignment(n_customers, frequencies, T):
    """
    Assign visit patterns to balance daily customer load.
    n_customers : number of customers
    frequencies : list of visit frequencies (one per customer)
    T           : number of days in planning horizon
    Returns list of assigned patterns (one per customer).
    """
    # Daily load tracker
    daily_load = [0] * T
    assigned = []
```

9.7 Conclusions and Future Research Directions I

This chapter has surveyed four important variants of the VRP, each capturing a different real-world feature that the classical model cannot handle.

Summary: What Each Variant Adds

Variant	Real-World Motivation	Key Constraint	Best Method (large n)
VRPB	Delivery + collection	Linehaul before backhaul	Tabu Search, ALNS
HFVRP	Mixed fleet (trucks of different sizes)	Choose vehicle type per route	Tabu Search, GA
PVRP	Weekly delivery schedules	Visit frequency + pattern assignment	Unified Tabu Search
SDVRP	Large demands, bin-packing waste	Full demand split across vehicles	B&P (exact), TS (heuristic)

9.7 Conclusions and Future Research Directions II

Open Research Questions

- 1 **Exact algorithms for large instances:** current exact methods (B&P) top out at ≈ 100 customers for all four variants. Advances in cut-generation, branch-and-cut-and-price, and decomposition are needed for larger instances.
- 2 **Hybrid variants:** all four variants can be combined (e.g. PVRP with heterogeneous fleet and backhauls). These combinations are increasingly common in practice but barely studied theoretically.
- 3 **Stochastic and dynamic extensions:** demands may be uncertain, customers may cancel, or new orders may arrive mid-route. Robust and online versions of VRPB, HFVRP, PVRP, and SDVRP are active research areas.
- 4 **Electric vehicles in HFVRP:** as fleets transition to electric vehicles (EVs), capacity constraints are joined by battery range constraints and charging decisions — a new dimension for the heterogeneous fleet problem.
- 5 **Machine learning for pattern selection (PVRP):** pattern assignment in PVRP is a combinatorial assignment problem that may benefit from reinforcement learning approaches that learn good assignment policies from historical routing data.

Appendix: VRPB Formulation Summary I

Appendix: VRPB Formulation Summary II

VRPB: Key Sets, Variables, and Constraints

Sets:

L = set of linehaul customers (delivery)

B = set of backhaul customers (pickup)

$V = \{0\} \cup L \cup B$, where 0 = depot

Variables:

$x_{ij}^k = 1$ if vehicle k travels arc (i, j) , else 0

$y_i^k = 1$ if customer i is served by vehicle k

Objective:

$\min \sum_k \sum_{(i,j)} c_{ij} x_{ij}^k$ (minimize total travel cost)

Key Constraints:


Capacity: $\sum_{i \in L} d_i y_i^k \leq Q$ (vehicle capacity not exceeded)

Precedence: linehaul visits before backhaul visits in each route

Each customer visited exactly once

Flow conservation at each node


Appendix: SDVRP Savings Bound Illustration I



Appendix: SDVRP Savings Bound Illustration II

This plot illustrates the theoretical bound on the ratio $VRP^*/SDVRP^*$ as a function of the demand fraction d_i/Q (the ratio of a customer's demand to the vehicle capacity). When demands are small relative to capacity (left side of the plot), the bound approaches 1 — splitting gives no benefit. When demands approach the full capacity (right side), the bound can be as large as 2 — standard VRP can cost up to twice as much as the optimal split-delivery solution. The shaded region represents the potential cost reduction available through split deliveries.

Appendix: VRPB Three Variants Side by Side I



Panel (a): The classical VRPB with strict precedence — all linehaul customers (blue) appear before all backhaul customers (red) in every route. Panel (b): The mixed VRPB (VRPMB) allows any interleaving of linehaul and backhaul customers. Panel (c): The clustered VRPB (VRPCB) groups customers of the same type in separate route segments.

Chapter 9: Summary Slide

VRPB Key Facts

- Two customer types: delivery (linehaul) and pickup (backhaul).
- Precedence: all deliveries before any pickup on each route.
- Three sub-variants: strict, mixed, clustered.
- Best heuristic: Tabu Search (Toth and Vigo, 1997).

PVRP Key Facts

- Planning horizon T periods; frequencies f_i per customer.
- Two sub-decisions: pattern assignment + daily routing.
- Benchmark: Cordeau et al. p-instances ($T = 5$).
- Best method: Unified Tabu Search (Cordeau et al., 1997).

HFVRP Key Facts

- T vehicle types with different capacities and costs.
- Objective: minimise fixed + variable costs.
- Free-fleet variant: also choose fleet composition.
- Benchmark: Golden et al. 12 instances.

SDVRP Key Facts

- Customer demand split across multiple vehicles.
- Lower bound on vehicles: $\lceil D/Q \rceil$.
- Savings: up to 50% cost reduction for large demands.
- Exact B&P handles ≤ 100 customers.